

Scientific Computing

Mathematical Foundations for Machine Learning

Course Conspect

Course Structure

Block I: Probability and Statistics

Lecture	Topic
1	Probability models, conditioning, and random variables
2	Distributions, expectation, and variance
3	Joint distributions, dependence, and random vectors
4	Samples, concentration, and limit theorems
5	Statistical estimation: method of moments, maximum likelihood, and MAP
6	Sampling distributions, confidence intervals, and bootstrap
7	Entropy, cross-entropy, KL divergence, Fisher information, and the Cramér–Rao bound

Block II: Linear Algebra, Calculus, and Optimization

Lecture	Topic
8	Geometry of data, projections, and least squares
9	Eigenvalues, positive-definite matrices, singular value decomposition, and PCA
10	Multivariable calculus, matrix calculus, and automatic differentiation
11	Optimization for machine learning

Block III: Numerical Computing and Autodiff

Lecture	Topic
12	NumPy, visualization, and numerical linear algebra
13	PyTorch, autograd, and the integrated capstone

Contents

Course Structure	i
1 Probability Models, Conditioning, and Random Variables	1
1.1 Probability spaces and events	1
1.2 Conditional probability and total probability	2
1.3 Independence and mutual exclusivity	3
1.4 Bayes' theorem and base rates	3
1.5 Random variables	4
1.6 PMFs, PDFs, and CDFs	5
2 Distributions, Expectation, and Variance	7
2.1 Exponential waiting times	7
2.2 Expectation and functions of a random variable	9
2.3 Variance and standard deviation	11
2.4 Poisson counts	13
2.5 Gamma waiting times and the event-count connection	15
2.6 Bernoulli and Binomial distributions	17
2.7 Geometric waiting times	19
2.8 The Gaussian distribution	20
2.9 Mixtures and heterogeneity	22
2.10 Extension: Gaussian scale mixtures and Student's t	27
2.11 Distribution reference	29

Chapter 1

Probability Models, Conditioning, and Random Variables

Probability starts with a model: a precise description of what may happen, which statements about the outcome are meaningful, and how probabilities are assigned. Conditioning changes the reference population, while a random variable turns outcomes into numerical values whose distribution can be studied.

1.1 Probability spaces and events

Definition 1.1 (Probability space). A probability space is a triple

$$(\Omega, \mathcal{F}, \mathbb{P}).$$

Here Ω is the *sample space*, whose elements are outcomes; $\mathcal{F}$ is a collection of subsets of Ω , called events; and $\mathbb{P}$ assigns a probability to every event.

The event collection $\mathcal{F}$ is closed under complements and countable unions:

$$A \in \mathcal{F} \implies A^c \in \mathcal{F}, \quad A_1, A_2, \dots \in \mathcal{F} \implies \bigcup_{i=1}^{\infty} A_i \in \mathcal{F}.$$

It also contains Ω . In a finite model one normally takes $\mathcal{F} = 2^\Omega$, the set of all subsets of Ω .

For events $A, B \subseteq \Omega$:

- $A \cap B$ means that both events occur;
- $A \cup B$ means that at least one occurs;
- $A^c = \Omega \setminus A$ means that A does not occur;
- $A \setminus B = A \cap B^c$ means that A occurs but B does not.

Definition 1.2 (Probability axioms). For events in $\mathcal{F}$, a probability measure satisfies:

$$\begin{aligned} \mathbb{P}(A) &\geq 0, \\ \mathbb{P}(\Omega) &= 1, \\ \mathbb{P}\left(\bigcup_{i=1}^{\infty} A_i\right) &= \sum_{i=1}^{\infty} \mathbb{P}(A_i) \quad \text{for pairwise disjoint } A_i. \end{aligned}$$

Several basic rules follow from these axioms:

$$\mathbb{P}(\emptyset) = 0, \quad (1.1)$$

$$\mathbb{P}(A^c) = 1 - \mathbb{P}(A), \quad (1.2)$$

$$A \subseteq B \implies \mathbb{P}(A) \leq \mathbb{P}(B), \quad (1.3)$$

$$\mathbb{P}(A \cup B) = \mathbb{P}(A) + \mathbb{P}(B) - \mathbb{P}(A \cap B). \quad (1.4)$$

The last identity is the inclusion–exclusion rule. Adding $\mathbb{P}(A)$ and $\mathbb{P}(B)$ counts the overlap twice, so it must be subtracted once.

Example 1.3 (Two fair coin tosses). Toss a fair coin twice and record the ordered result. Then

$$\Omega = \{HH, HT, TH, TT\}, \quad \mathbb{P}(\{\omega\}) = \frac{1}{4} \quad \text{for every } \omega \in \Omega.$$

Let

$$A = \{\text{first toss is } H\} = \{HH, HT\}, \quad B = \{\text{second toss is } H\} = \{HH, TH\}.$$

Then

$$A \cap B = \{HH\}, \quad A \cup B = \{HH, HT, TH\}, \quad A^c = \{TH, TT\}.$$

Inclusion–exclusion gives

$$\mathbb{P}(A \cup B) = \frac{1}{2} + \frac{1}{2} - \frac{1}{4} = \frac{3}{4}.$$

For a finite sample space with equally likely outcomes,

$$\mathbb{P}(A) = \frac{|A|}{|\Omega|}.$$

This is a counting shortcut, not an additional axiom: it is valid only after equal likelihood has been justified.

1.2 Conditional probability and total probability

Definition 1.4 (Conditional probability). For $\mathbb{P}(B) > 0$, the probability of A given B is

$$\mathbb{P}(A \mid B) = \frac{\mathbb{P}(A \cap B)}{\mathbb{P}(B)}. \quad (1.5)$$

Conditioning restricts the reference set to B and renormalizes the remaining probabilities. The event after the vertical bar determines the denominator. In general, $\mathbb{P}(A \mid B)$ and $\mathbb{P}(B \mid A)$ are different quantities.

Rearranging (1.5) gives the multiplication rule

$$\mathbb{P}(A \cap B) = \mathbb{P}(A \mid B)\mathbb{P}(B) = \mathbb{P}(B \mid A)\mathbb{P}(A). \quad (1.6)$$

Example 1.5 (At least one head). In the two-toss model, let

$$E = \{\text{at least one head}\} = \{HH, HT, TH\}.$$

Given E , the probability that the second toss is heads is

$$\mathbb{P}(B \mid E) = \frac{\mathbb{P}(B \cap E)}{\mathbb{P}(E)} = \frac{2/4}{3/4} = \frac{2}{3}.$$

The unconditional probability is $\mathbb{P}(B) = 1/2$. Conditioning removes the outcome TT , in which B fails.

Theorem 1.6 (Law of total probability). *Let $B_1, \dots, B_m$ be a partition of Ω : the events are pairwise disjoint and their union is Ω . If $\mathbb{P}(B_i) > 0$, then*

$$\mathbb{P}(A) = \sum_{i=1}^m \mathbb{P}(A \mid B_i) \mathbb{P}(B_i). \quad (1.7)$$

Indeed, the events $A \cap B_i$ are disjoint and their union is A . Add their probabilities and then apply the multiplication rule. For a binary partition B, B^c ,

$$\mathbb{P}(A) = \mathbb{P}(A \mid B) \mathbb{P}(B) + \mathbb{P}(A \mid B^c) \mathbb{P}(B^c).$$

1.3 Independence and mutual exclusivity

Definition 1.7 (Independence). Events A and B are independent if

$$\mathbb{P}(A \cap B) = \mathbb{P}(A) \mathbb{P}(B). \quad (1.8)$$

If $\mathbb{P}(B) > 0$, this is equivalent to $\mathbb{P}(A \mid B) = \mathbb{P}(A)$.

Independence means that learning whether one event occurred does not change the probability of the other. In the coin example,

$$\mathbb{P}(A \cap B) = \frac{1}{4} = \frac{1}{2} \cdot \frac{1}{2},$$

so the events “first toss is heads” and “second toss is heads” are independent.

Definition 1.8 (Mutual exclusivity). Events A and B are mutually exclusive, or disjoint, if

$$A \cap B = \emptyset.$$

They cannot occur together.

Proposition 1.9. *Two mutually exclusive events with positive probability are not independent.*

Proof. Mutual exclusivity gives $\mathbb{P}(A \cap B) = 0$, while positive probabilities give $\mathbb{P}(A) \mathbb{P}(B) > 0$. Thus (1.8) cannot hold. $\square$

The distinction is essential:

	Mutual exclusivity	Independence
Can both occur?	No	Usually yes
Defining condition	$A \cap B = \emptyset$	$\mathbb{P}(A \cap B) = \mathbb{P}(A) \mathbb{P}(B)$
Information meaning	One rules out the other	One does not change the other

1.4 Bayes' theorem and base rates

Theorem 1.10 (Bayes' theorem). *Let $B_1, \dots, B_m$ be a partition of Ω . For $\mathbb{P}(A) > 0$,*

$$\mathbb{P}(B_j \mid A) = \frac{\mathbb{P}(A \mid B_j) \mathbb{P}(B_j)}{\sum_{i=1}^m \mathbb{P}(A \mid B_i) \mathbb{P}(B_i)}. \quad (1.9)$$

The numerator is the probability of the route B_j followed by A . The denominator adds every mutually exclusive route that can produce A .

Example 1.11 (Rare-defect classifier). Let D mean that an item is defective and $+$ mean that a classifier produces a positive flag. Assume

$$\mathbb{P}(D) = 0.01, \quad \mathbb{P}(+ | D) = 0.99, \quad \mathbb{P}(+ | D^c) = 0.05.$$

The probability of a positive result is

$$\mathbb{P}(+) = 0.99(0.01) + 0.05(0.99) = 0.0594.$$

Therefore

$$\mathbb{P}(D | +) = \frac{0.99(0.01)}{0.0594} = \frac{1}{6} \approx 0.167.$$

The same calculation can be expressed with expected counts in a population of 10,000 items:

Actual state	Positive	Negative	Total
Defect D	99	1	100
No defect D^c	495	9,405	9,900
Total	594	9,406	10,000

Only 99 of the 594 positive items are defective. The sensitivity $\mathbb{P}(+ | D) = 0.99$ is not the posterior $\mathbb{P}(D | +)$. The low base rate and the false positives among the much larger non-defective group dominate the denominator.

For a Bayesian classifier with classes $C_1, \dots, C_K$ and observed features x ,

$$\mathbb{P}(C_k | x) = \frac{p(x | C_k)\mathbb{P}(C_k)}{\sum_{j=1}^K p(x | C_j)\mathbb{P}(C_j)}. \quad (1.10)$$

The factors have distinct roles:

- $\mathbb{P}(C_k)$ is the prior or class base rate;
- $p(x | C_k)$ is the class-conditional likelihood;
- $\mathbb{P}(C_k | x)$ is the posterior after observing x .

The maximum a posteriori rule is

$$\hat{C}_{\text{MAP}}(x) = \arg \max_k p(x | C_k)\mathbb{P}(C_k).$$

1.5 Random variables

Definition 1.12 (Random variable). A real-valued random variable is a function

$$X : \Omega \longrightarrow \mathbb{R}$$

such that $\{\omega : X(\omega) \leq x\}$ is an event for every $x \in \mathbb{R}$. In finite examples this event condition is automatic.

The sampled outcome is random; the mapping X is fixed.

Example 1.13 (Number of heads). For two coin tosses, let X be the number of heads. Then

Outcome ω	HH	HT	TH	TT
$X(\omega)$	2	1	1	0

The support $\{0, 1, 2\}$ is the set of possible values of X ; it is not the original sample space.

A *discrete* random variable takes values in a finite or countable set. An *absolutely continuous* random variable assigns probabilities to intervals through a density. Mixed random variables, with both point masses and a continuous part, also exist.

1.6 PMFs, PDFs, and CDFs

1.6.1 Probability mass functions

Definition 1.14 (PMF). For a discrete random variable X , its probability mass function is

$$p_X(x) = \mathbb{P}(X = x), \quad p_X(x) \geq 0, \quad \sum_x p_X(x) = 1.$$

For the number of heads in two fair tosses,

$$p_X(0) = \frac{1}{4}, \quad p_X(1) = \frac{1}{2}, \quad p_X(2) = \frac{1}{4}.$$

The probability of an event is obtained by summing the masses of its values; for example, $\mathbb{P}(X \geq 1) = p_X(1) + p_X(2) = 3/4$.

1.6.2 Cumulative distribution functions

Definition 1.15 (CDF). For every real-valued random variable X , the cumulative distribution function is

$$F_X(x) = \mathbb{P}(X \leq x).$$

Every CDF is nondecreasing and right-continuous, with

$$\lim_{x \rightarrow -\infty} F_X(x) = 0, \quad \lim_{x \rightarrow \infty} F_X(x) = 1.$$

It determines interval probabilities through

$$\mathbb{P}(a < X \leq b) = F_X(b) - F_X(a). \quad (1.11)$$

For the number of heads,

$$F_X(x) = \begin{cases} 0, & x < 0, \\ \frac{1}{4}, & 0 \leq x < 1, \\ \frac{3}{4}, & 1 \leq x < 2, \\ 1, & x \geq 2. \end{cases}$$

For a discrete variable, the jump at x equals its point mass:

$$p_X(x) = F_X(x) - F_X(x^-),$$

where $F_X(x^-)$ denotes the left limit.

1.6.3 Probability density functions

Definition 1.16 (PDF). An absolutely continuous random variable Y has a probability density function f_Y satisfying

$$f_Y(y) \geq 0, \quad \int_{-\infty}^{\infty} f_Y(y) dy = 1,$$

and

$$\mathbb{P}(a < Y \leq b) = \int_a^b f_Y(y) dy. \quad (1.12)$$

Warning. A density value is not a point probability. For an absolutely continuous random variable, $\mathbb{P}(Y = y) = 0$ for every single point, and $f_Y(y)$ may be larger than one. Probability is the area under the density.

The CDF and density are related by

$$F_Y(y) = \int_{-\infty}^y f_Y(t) dt, \quad F'_Y(y) = f_Y(y)$$

at points where the derivative exists.

Example 1.17 (Uniform distribution on the unit interval). Let

$$f_Y(y) = \begin{cases} 1, & 0 \leq y \leq 1, \\ 0, & \text{otherwise.} \end{cases}$$

Then

$$F_Y(y) = \begin{cases} 0, & y < 0, \\ y, & 0 \leq y \leq 1, \\ 1, & y > 1. \end{cases}$$

Thus

$$\mathbb{P}(0.2 < Y \leq 0.7) = \int_{0.2}^{0.7} 1 dy = F_Y(0.7) - F_Y(0.2) = 0.5.$$

Chapter 2

Distributions, Expectation, and Variance

2.1 Exponential waiting times

Imagine watching a help desk's inbox. Start a stopwatch now and stop it when the next message arrives. The number on the stopwatch is T , the waiting time in minutes. On one occasion we might wait 0.3 minutes; on another, 7 minutes. We want to answer questions such as: what is the probability that we are still waiting after five minutes?

2.1.1 A model of arrivals without a schedule

First build a simple model using one-second intervals. Suppose that in each interval a message arrives with probability $1/300$, independently of what happened in earlier intervals. Most seconds contain no message. Occasionally two successful seconds occur close together; occasionally there is a long gap.

In this model, waiting more than five minutes means getting no message in any of 300 consecutive seconds. By independence,

$$\mathbb{P}(T > 5) = \left(1 - \frac{1}{300}\right)^{300} \approx 0.3673.$$

A long wait is therefore quite plausible. The model does not place an arrival at regular five-minute intervals.

Now shorten the time intervals. Write λ for the arrival rate, measured in messages per minute. A small interval of length Δ minutes has arrival probability $\lambda\Delta$. In our example, $\lambda = 0.2$, because $0.2 \times (1/60) = 1/300$. The rate describes how frequently messages arrive on average; it is not itself a probability.

Divide a time span of t minutes into n equal intervals. In the discrete approximation, each has arrival probability $\lambda t/n$. The probability that all of them are empty is

$$\left(1 - \frac{\lambda t}{n}\right)^n.$$

Letting the intervals become arbitrarily short gives

$$\mathbb{P}(T > t) = \lim_{n \rightarrow \infty} \left(1 - \frac{\lambda t}{n}\right)^n = e^{-\lambda t}, \quad t \geq 0. \quad (2.1)$$

This is the probability that *no message has arrived by time t* . The exponential function appears because we multiply the no-arrival probabilities of many short intervals. The resulting continuous waiting time has an *exponential distribution*.

The assumptions describe an inbox with no timetable and no change in activity during the period being studied. They would be unsuitable for notifications sent every hour or an inbox that becomes much busier at lunchtime.

2.1.2 From the waiting probability to the density

A message has either arrived by time t , or it has not. Thus the CDF is

$$F_T(t) = \mathbb{P}(T \leq t) = 1 - e^{-\lambda t}, \quad t \geq 0.$$

Differentiating this CDF, using the relation from Chapter 1, gives the density of the waiting time.

Definition 2.1 (Exponential distribution). For $\lambda > 0$, the notation $T \sim \text{Exp}(\lambda)$ means

$$f_T(t) = \begin{cases} \lambda e^{-\lambda t}, & t \geq 0, \\ 0, & t < 0. \end{cases}$$

The parameter λ is the arrival rate.

Increasing λ makes short waits more likely. For example, at rate 0.2 per minute the probability of waiting more than five minutes is $e^{-1} \approx 0.3679$. At rate 0.5, that probability drops to $e^{-2.5} \approx 0.0821$. Figure 2.1 shows this comparison directly.




Figure 2.1: The probability that the inbox is still empty after t minutes. Every curve starts at one: at time zero we have not yet waited. A busier inbox has a larger rate and a faster decrease in the probability of still waiting. The dashed vertical line marks five minutes.

For an interval question, subtract two waiting probabilities. With $\lambda = 0.2$,

$$\mathbb{P}(2 < T \leq 5) = \mathbb{P}(T > 2) - \mathbb{P}(T > 5) = e^{-0.4} - e^{-1} \approx 0.3024.$$

We keep the cases that are still waiting at minute two and remove those that are still waiting at minute five.

A different practical question is how long we must wait to have seen a message in 95% of repetitions. Solve

$$1 - e^{-0.2t} = 0.95 \implies t = -\frac{\log(0.05)}{0.2} \approx 14.98 \text{ minutes.}$$

This cutoff is the 95-th percentile of the waiting time. More generally, the q -quantile is $t_q = -\log(1 - q)/\lambda$, for $0 < q < 1$.

2.1.3 Does a long wait make an arrival overdue?

Suppose we have already waited ten minutes with no message. In the model, the next short interval has the same arrival chance as it did at the start. Earlier empty intervals do not change that chance. Therefore the probability of waiting another five minutes is still $e^{-0.2 \cdot 5} \approx 0.3679$.

We can verify this using conditional probability. For $s, t \geq 0$,

$$\mathbb{P}(T > s + t \mid T > s) = \frac{\mathbb{P}(T > s + t)}{\mathbb{P}(T > s)} = \frac{e^{-\lambda(s+t)}}{e^{-\lambda s}} = e^{-\lambda t}.$$

The right-hand side is $\mathbb{P}(T > t)$. The distribution of the remaining wait is the same as the distribution of the original wait. This property is called *memorylessness*.

A scheduled bus behaves differently: if it comes every ten minutes, waiting nine minutes tells us something about the remaining wait. The exponential model is appropriate only when this kind of schedule, ageing, or change in activity is absent.

2.2 Expectation and functions of a random variable

The probability of a long wait answers one question. The average amount of waiting answers another. If we repeat the inbox experiment independently many times and average the recorded times, what value should that average be close to?

Start with a simpler example. Suppose a task takes one minute with probability $3/4$ and nine minutes with probability $1/4$. In many repetitions, about three quarters of the tasks take one minute and one quarter take nine minutes. Their average is close to

$$\frac{3}{4} \cdot 1 + \frac{1}{4} \cdot 9 = 3 \text{ minutes.}$$

This probability-weighted average is the *expectation*. It can be three minutes even though no individual task takes three minutes.

2.2.1 Averaging with probabilities

For a discrete random variable, weight each possible value by its probability and add. For a continuous variable, a small interval near x has probability approximately $f_X(x)\Delta x$, so its contribution to the average is approximately $xf_X(x)\Delta x$. Summing these contributions over increasingly small intervals gives an integral.

Definition 2.2 (Expectation). The expectation of X is

$$\mathbb{E}[X] = \begin{cases} \sum x p_X(x), & X \text{ discrete,} \\ \int_{-\infty}^{\infty} x f_X(x) dx, & X \text{ has a density.} \end{cases}$$

For a finite expectation we require $\mathbb{E}[|X|] < \infty$: the sum or integral of absolute contributions must converge.

This condition matters when very large values, although rare, contribute so much that no finite average exists. All the distributions in the main part of this chapter have finite means. The Student extension provides an example where the condition can fail.

For exponential waiting times, the definition gives

$$\mathbb{E}[T] = \int_0^\infty t \lambda e^{-\lambda t} dt.$$

To evaluate it, integrate by parts with $u = t$ and $dv = \lambda e^{-\lambda t} dt$:

$$\mathbb{E}[T] = [-te^{-\lambda t}]_0^\infty + \int_0^\infty e^{-\lambda t} dt = \frac{1}{\lambda}.$$

At rate 0.2 messages per minute, the mean wait is five minutes. This is an average over repeated waits, not an appointment for the next arrival. We already found that more than a third of waits exceed five minutes.

2.2.2 First transform the observation, then average

Sometimes we want to average something other than the waiting time itself. For example, assign a response score

$$g(t) = e^{-at}, \quad a > 0.$$

An immediate response scores one; longer waits receive smaller scores. The parameter a controls how quickly the score decreases with time.

To find the average score, calculate the score for each possible wait and then weight it by the probability of that wait:

$$\mathbb{E}[g(T)] = \int_0^\infty g(t) f_T(t) dt.$$

For this particular score,

$$\mathbb{E}[e^{-aT}] = \int_0^\infty e^{-at} \lambda e^{-\lambda t} dt = \frac{\lambda}{\lambda + a}.$$

This calculation illustrates the general rule

$$\mathbb{E}[g(X)] = \begin{cases} \sum g(x) p_X(x), & X \text{ discrete,} \\ \int_{-\infty}^{\infty} g(x) f_X(x) dx, & X \text{ has a density.} \end{cases} \quad (2.2)$$

We assume $\mathbb{E}[|g(X)|] < \infty$. The rule lets us use the distribution of X directly, without first deriving the distribution of $g(X)$.

Scoring the average wait is a different operation. For $a = \lambda = 1$,

$$\mathbb{E}[e^{-T}] = \frac{1}{2}, \quad e^{-\mathbb{E}[T]} = e^{-1} \approx 0.3679.$$

The nonlinear transformation makes the order of the two operations matter.

Another useful choice is $g(t) = t^2$. Large waits receive much more weight after squaring. Again integrating by parts,

$$\mathbb{E}[T^2] = \int_0^\infty t^2 \lambda e^{-\lambda t} dt = \frac{2}{\lambda} \mathbb{E}[T] = \frac{2}{\lambda^2}.$$

We will use this second moment to measure spread.

2.2.3 Why expectations add

Suppose completing a job requires a wait X and a processing time Y . On each repetition, the total is $X + Y$. Averaging the totals gives the average waiting time plus the average processing time. This remains true even if long waits tend to accompany long processing times.

Theorem 2.3 (Linearity of expectation). *Let $X_1, \dots, X_n$ have finite expectations. For constants a_i and b ,*

$$\mathbb{E}\left[b + \sum_{i=1}^n a_i X_i\right] = b + \sum_{i=1}^n a_i \mathbb{E}[X_i].$$

Independence is not required.

For a finite sample space, distributing the sum makes the reason explicit:

$$\sum_{\omega \in \Omega} (X(\omega) + Y(\omega)) \mathbb{P}(\{\omega\}) = \mathbb{E}[X] + \mathbb{E}[Y].$$

The integral version works in the same way.

For instance, if a job has a fixed two-minute setup followed by $T \sim \text{Exp}(0.2)$, its mean duration is $2 + \mathbb{E}[T] = 7$ minutes. Expressing T in seconds instead multiplies its mean by 60.

Linearity also makes counts easy to analyze. For an event A , the indicator $\mathbf{1}_A$ is one if A occurs and zero otherwise. Its expectation is

$$\mathbb{E}[\mathbf{1}_A] = 1 \cdot \mathbb{P}(A) + 0 \cdot \mathbb{P}(A^c) = \mathbb{P}(A).$$

Consequently, the expected number of events occurring among $A_1, \dots, A_n$ is

$$\mathbb{E}\left[\sum_{i=1}^n \mathbf{1}_{A_i}\right] = \sum_{i=1}^n \mathbb{P}(A_i).$$

No independence assumption is needed for this mean. Dependence will matter when we ask about the count's spread.

2.3 Variance and standard deviation

A constant five-minute wait and an exponential wait with mean five minutes have the same expectation. Their behavior is very different: the first is perfectly predictable, while the second sometimes ends almost immediately and sometimes takes much longer. We need a measure of this variation around the mean.

If $\mu = \mathbb{E}[X]$, the deviation of an observation is $X - \mu$. Averaging the signed deviations gives zero: $\mathbb{E}[X - \mu] = 0$. Positive and negative deviations cancel. Squaring them prevents that cancellation and gives more weight to large departures from the mean.

Definition 2.4 (Variance and standard deviation). For X with a finite second moment,

$$\text{Var}(X) = \mathbb{E}[(X - \mu)^2], \quad \mu = \mathbb{E}[X], \quad \text{sd}(X) = \sqrt{\text{Var}(X)}.$$

Variance is the average squared deviation. If waiting time is in minutes, variance is in minutes squared. Taking the square root returns to minutes; this is why standard deviation is useful when describing the size of typical fluctuations. It is not a bound on individual observations.

Expanding the square gives a convenient way to calculate variance:

$$\begin{aligned}\text{Var}(X) &= \mathbb{E}[X^2 - 2\mu X + \mu^2] \\ &= \mathbb{E}[X^2] - 2\mu\mathbb{E}[X] + \mu^2 = \mathbb{E}[X^2] - \mathbb{E}[X]^2.\end{aligned}\tag{2.3}$$

For exponential waiting times, we have already computed both terms:

$$\text{Var}(T) = \frac{2}{\lambda^2} - \frac{1}{\lambda^2} = \frac{1}{\lambda^2}, \quad \text{sd}(T) = \frac{1}{\lambda}.$$

The inbox with rate 0.2 has both mean and standard deviation equal to five minutes. A constant five-minute wait instead has variance zero.

2.3.1 Changing the origin or the units

Adding a fixed delay b moves every observation and its mean by the same amount. The deviations from the mean stay unchanged. Multiplying times by a multiplies every deviation by a , so their squares are multiplied by a^2 . Therefore

$$\mathbb{E}[aX + b] = a\mathbb{E}[X] + b, \quad \text{Var}(aX + b) = a^2 \text{Var}(X).$$

For example, converting minutes to seconds multiplies the mean and standard deviation by 60, and the variance by 60^2 . Variance is not linear.

2.3.2 Adding two variable delays

For a job with waiting time X and processing time Y , let $U = X - \mathbb{E}[X]$ and $V = Y - \mathbb{E}[Y]$. The deviation of the total duration is $U + V$. Squaring gives

$$(U + V)^2 = U^2 + V^2 + 2UV.$$

The first two terms describe the separate spreads. The last term tells us whether the two delays fluctuate together. If both are above their means, or both are below, their product UV is positive. If one is above its mean while the other is below, the product is negative.

The average of this product is called the *covariance*:

$$\text{Cov}(X, Y) = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])].$$

Taking expectations in the squared-sum identity gives

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y) + 2 \text{Cov}(X, Y).\tag{2.4}$$

For independent variables, expectations of products factor. Thus $\mathbb{E}[UV] = \mathbb{E}[U]\mathbb{E}[V] = 0$, and the covariance term vanishes. For independent variables with finite second moments, this extends to

$$\text{Var}\left(\sum_{i=1}^n a_i X_i\right) = \sum_{i=1}^n a_i^2 \text{Var}(X_i).$$

Example 2.5 (Three different ways to combine fluctuations). Let X, Y each have variance one. If they are independent, the variance of $X + Y$ is two. If $Y = X$, the fluctuations reinforce each other and $\text{Var}(X + Y) = \text{Var}(2X) = 4$. If $Y = -X$, they cancel exactly and the sum has variance zero.

Zero covariance is enough for variances to add, even without full independence. We will study that distinction further when considering joint distributions. For the sums used below, independence will be an explicit part of the model.

2.4 Poisson counts

Return to the inbox, but change the question. Instead of timing the next message, watch for ten minutes and count how many messages arrive. A waiting time can be 2.4 minutes; a message count must be an integer such as zero, one, or two. The same arrivals produce two different kinds of random variables.

Write $N(t)$ for the number of messages arriving between time zero and time t . Figure 2.2 shows one possible record. At minute six we have counted two messages; at minute ten we have counted three.




Figure 2.2: One possible sequence of arrivals, at 1.5, 4, and 8 minutes. The gaps W_1, W_2, W_3 are waiting times; the arrival times T_k are cumulative sums of those gaps. By the marked time $t = 6$, two messages have arrived, so $N(6) = 2$.

2.4.1 Counting successful short intervals

We can use the same short-interval model as before. Divide t minutes into n tiny intervals, each independently containing a message with probability $p_n = \lambda t/n$. To get exactly k messages, exactly k intervals must contain a message and the others must be empty.

For a particular choice of those k intervals, independence gives probability $p_n^k (1 - p_n)^{n-k}$. The number of ways to choose the intervals is

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}.$$

For example, if there are three intervals and exactly one message, it can be in the first, second, or third interval: three possibilities. Here $n! = 1 \cdot 2 \cdots n$, with $0! = 1$. Adding the equally probable possibilities gives

$$\binom{n}{k} \left(\frac{\lambda t}{n}\right)^k \left(1 - \frac{\lambda t}{n}\right)^{n-k}.$$

Put $m = \lambda t$. For fixed k , as n increases,

$$\frac{\binom{n}{k}}{n^k} \longrightarrow \frac{1}{k!}, \quad \left(1 - \frac{m}{n}\right)^{n-k} \longrightarrow e^{-m}.$$

Thus the continuous-time count has probability $e^{-m} m^k / k!$ of equalling k . This is the *Poisson distribution*.

Definition 2.6 (Poisson distribution). For $m > 0$, a count $N \sim \text{Poisson}(m)$ has PMF

$$\mathbb{P}(N = k) = e^{-m} \frac{m^k}{k!}, \quad k = 0, 1, 2, \dots$$

For our inbox observed for t minutes, $N(t) \sim \text{Poisson}(\lambda t)$.

The rate λ describes arrivals per minute; $m = \lambda t$ describes the expected number over the whole observation window. At rate 0.2 per minute, a ten-minute window has $m = 2$, while a thirty-minute window has $m = 6$. Neither count is guaranteed to equal its mean.

Example 2.7 (No messages or several messages). For the ten-minute count with $m = 2$,

$$\mathbb{P}(N(10) = 0) = e^{-2} \approx 0.1353,$$

and

$$\mathbb{P}(N(10) \geq 2) = 1 - \mathbb{P}(N(10) = 0) - \mathbb{P}(N(10) = 1) = 1 - e^{-2}(1 + 2) \approx 0.5940.$$

No messages in ten minutes is exactly the event that the first waiting time exceeds ten minutes. The Poisson and Exponential calculations agree: $\mathbb{P}(N(10) = 0) = \mathbb{P}(T > 10) = e^{-2}$.

2.4.2 Mean and variance of the count

To check normalization, sum the PMF and use the exponential series:

$$\sum_{k=0}^{\infty} e^{-m} \frac{m^k}{k!} = e^{-m} e^m = 1.$$

For the mean, the factor k cancels one factor in $k!$. Reindexing the remaining sum with $j = k - 1$ gives

$$\mathbb{E}[N] = \sum_{k=1}^{\infty} k e^{-m} \frac{m^k}{k!} = m \sum_{j=0}^{\infty} e^{-m} \frac{m^j}{j!} = m.$$

For the variance we need $\mathbb{E}[N^2]$. It is easier to first calculate $\mathbb{E}[N(N - 1)]$, since $k(k - 1)$ cancels two factors in $k!$:

$$\mathbb{E}[N(N - 1)] = m^2 \sum_{j=0}^{\infty} e^{-m} \frac{m^j}{j!} = m^2.$$

Now use $N^2 = N(N - 1) + N$:

$$\mathbb{E}[N^2] = m^2 + m, \quad \text{Var}(N) = m.$$

For ten minutes, the mean and variance are both two, and the standard deviation is $\sqrt{2}$ messages.

The equality of mean and variance will later help us recognize when a single Poisson model misses part of the variation in a population.

2.4.3 Combining independent arrival streams

Suppose two independent inboxes receive messages at rates λ_1 and λ_2 . Counting all their messages over the same interval should add the expected counts. In fact, the entire distribution remains Poisson:

$$N_1 + N_2 \sim \text{Poisson}(m_1 + m_2), \quad m_i = \lambda_i t.$$

To see why, split a total of k messages into j from the first inbox and $k - j$ from the second. Independence lets us multiply their probabilities; summing over j and using the binomial theorem gives

$$\mathbb{P}(N_1 + N_2 = k) = \frac{e^{-(m_1+m_2)}}{k!} \sum_{j=0}^k \binom{k}{j} m_1^j m_2^{k-j} = \frac{e^{-(m_1+m_2)} (m_1 + m_2)^k}{k!}.$$

In the short-interval model, counts in disjoint time windows are also independent. If arrivals instead trigger further arrivals or the rate changes with the time of day, these assumptions need to be reconsidered.

2.5 Gamma waiting times and the event-count connection

Suppose the help desk collects three messages before starting a batch. Now we need the time until the third message, rather than the time until the next one or the number received by a fixed deadline.

Let W_1 be the wait from time zero to the first message, W_2 the gap from the first to the second, and so on. In our arrival model these gaps are independent and each follows $\text{Exp}(\lambda)$. The k -th arrival time is

$$T_k = W_1 + \cdots + W_k.$$

For example, the gaps in Figure 2.2 are 1.5, 2.5, and 4 minutes, so $T_3 = 8$ minutes.

Before deriving a new density, we can already find the mean and variance:

$$\mathbb{E}[T_k] = \frac{k}{\lambda}, \quad \text{Var}(T_k) = \frac{k}{\lambda^2}.$$

Expectation adds by linearity; variance adds because the gaps are independent. At rate 0.2, waiting for three messages takes 15 minutes on average, with variance 75 minutes squared.

2.5.1 The distribution of a sum of waiting times

Consider two messages first. A total wait near five minutes could consist of a first gap near one minute and a second near four, or a first near two and a second near three, and so on. To find the density of the total, we must account for every possible split.

If the first gap is u , the second must be near $t - u$. Independence gives a product of densities for these two contributions. Integrating over $0 < u < t$ adds all the possible splits:

$$\begin{aligned} f_{T_2}(t) &= \int_0^t f_{W_1}(u) f_{W_2}(t-u) du \\ &= \int_0^t \lambda e^{-\lambda u} \lambda e^{-\lambda(t-u)} du = \lambda^2 t e^{-\lambda t}, \quad t > 0. \end{aligned}$$

This way of finding the density of a sum is called *convolution*.

Repeating the same step for k independent gaps gives

$$f_{T_k}(t) = \frac{\lambda^k}{(k-1)!} t^{k-1} e^{-\lambda t}, \quad t > 0.$$

For $k = 1$, this is the exponential density. For $k > 1$, the density starts near zero: fitting several arrivals into a very short time is unlikely. As t increases, enough time becomes available for all k arrivals. Eventually the density decreases because exceptionally long total waits are also unlikely.




Figure 2.3: Time until the first, second, or fourth arrival, at rate one. One arrival can occur almost immediately. Several arrivals require a sum of waiting times, moving the typical total wait to the right. The means are 1, 2, 4; the standard deviations are 1, $\sqrt{2}$, 2.

2.5.2 From an integer number of arrivals to a shape parameter

The density above gives a useful family of positive, often skewed distributions. We can extend its shape beyond integer values of k . The only obstacle is the factorial $(k-1)!$. Its continuous extension is the gamma function:

$$\Gamma(\alpha) = \int_0^\infty u^{\alpha-1} e^{-u} du, \quad \alpha > 0.$$

Integration by parts gives $\Gamma(\alpha+1) = \alpha\Gamma(\alpha)$. Since $\Gamma(1) = 1$, this implies $\Gamma(k) = (k-1)!$ for positive integers k .

Replacing k by a positive real number α and the factorial by $\Gamma(\alpha)$ gives the *Gamma distribution*. We write its rate as β so that λ can continue to denote the particular inbox rate.

Definition 2.8 (Gamma distribution). For shape $\alpha > 0$ and rate $\beta > 0$, $X \sim \text{Gamma}(\alpha, \beta)$ has density

$$f_X(x) = \frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}, \quad x > 0,$$

and zero density for $x < 0$.

The substitution $u = \beta x$ shows that its integral is one. The case $\alpha = 1$ is Exponential; integer $\alpha = k$ describes waiting for k arrivals at rate β . For noninteger shape, the density still makes sense, but there is no literal fractional-number-of-arrivals interpretation.

We can calculate the moments for every positive shape by inserting x or x^2 into the density integral:

$$\begin{aligned} \mathbb{E}[X] &= \frac{\Gamma(\alpha+1)}{\beta\Gamma(\alpha)} = \frac{\alpha}{\beta}, \\ \mathbb{E}[X^2] &= \frac{\Gamma(\alpha+2)}{\beta^2\Gamma(\alpha)} = \frac{\alpha(\alpha+1)}{\beta^2}. \end{aligned}$$

Subtracting the squared mean gives

$$\text{Var}(X) = \frac{\alpha}{\beta^2}.$$

Increasing the rate β shortens the waits. Increasing the shape α at fixed rate increases the mean, while the spread relative to that mean decreases: $\text{sd}(X)/\mathbb{E}[X] = 1/\sqrt{\alpha}$.

Throughout these notes the second Gamma parameter is a rate. Another common convention uses a scale $\theta = 1/\beta$, measured in the same units as X . In that convention the mean and variance are $\alpha\theta$ and $\alpha\theta^2$. Checking which parameterization is used prevents a reciprocal error.

2.5.3 A deadline question can be answered by counting

To have received the third message by minute ten means to have counted at least three messages by minute ten. These are the same event, expressed using different random variables:

$$\{T_k \leq t\} = \{N(t) \geq k\}, \quad k \geq 1. \quad (2.5)$$

Therefore, instead of integrating a Gamma density, we can add Poisson probabilities:

$$\mathbb{P}(T_k \leq t) = 1 - \mathbb{P}(N(t) < k) = 1 - e^{-\lambda t} \sum_{j=0}^{k-1} \frac{(\lambda t)^j}{j!}.$$

For rate 0.2, a ten-minute window has expected count two, so

$$\mathbb{P}(T_3 \leq 10) = 1 - e^{-2}(1 + 2 + 2^2/2) \approx 0.3233.$$

Receiving exactly k messages also has a waiting-time description: the k -th has arrived, but the $(k+1)$ -th has not. For $k \geq 1$, $\mathbb{P}(N(t) = k) = \mathbb{P}(T_k \leq t) - \mathbb{P}(T_{k+1} \leq t)$. For $k = 0$, it is simply $\mathbb{P}(T_1 > t)$.

2.6 Bernoulli and Binomial distributions

The inbox model asked whether a message arrived during each small time interval. Many other observations are also binary: a link was clicked or it was not, a transmission succeeded or failed, an item passed an inspection or failed it. We now describe these trials without reference to time.

2.6.1 One binary outcome

Code a click as one and no click as zero. If the click probability is p , then

$$\mathbb{P}(B = 1) = p, \quad \mathbb{P}(B = 0) = 1 - p, \quad 0 \leq p \leq 1.$$

This is a *Bernoulli distribution*, written $B \sim \text{Bernoulli}(p)$. The same construction applies to the indicator of any event.

For a binary variable, squaring changes neither possible value: $0^2 = 0$, $1^2 = 1$. Thus

$$\mathbb{E}[B] = p, \quad \mathbb{E}[B^2] = p, \quad \text{Var}(B) = p(1 - p).$$

At $p = 1/2$, neither outcome is favored and the variance is largest. At $p = 0$ or $p = 1$, every observation is the same and variance is zero.

2.6.2 A fixed number of independent trials

Show a link to n people. Assume their clicks are independent and each has the same probability p . If B_i indicates the i -th click, then the total number of clicks is

$$S_n = B_1 + \cdots + B_n.$$

It follows a *Binomial distribution*, $S_n \sim \text{Binomial}(n, p)$, with

$$\mathbb{P}(S_n = k) = \binom{n}{k} p^k (1-p)^{n-k}, \quad k = 0, \dots, n.$$

The counting argument is the one used for the short arrival intervals: choose the k successful trials, multiply their success and failure probabilities, and add over all choices. The displayed expression is read directly for $0 < p < 1$; at $p = 0$ or $p = 1$ the count is fixed.

We do not need to sum this PMF to obtain its moments. Since the count is a sum of independent Bernoulli variables,

$$\mathbb{E}[S_n] = np, \quad \text{Var}(S_n) = np(1-p).$$

For $n \geq 1$, the observed click proportion is $\hat{p} = S_n/n$. The scaling rules for mean and variance give

$$\mathbb{E}[\hat{p}] = p, \quad \text{Var}(\hat{p}) = \frac{p(1-p)}{n}.$$

With more independent observations, the proportion fluctuates less. For example, multiplying n by four halves its standard deviation.

Example 2.9 (A small click probability). For 100 independent impressions with $p = 0.03$, the count has mean 3 and variance 2.91. The chance of no clicks is

$$\mathbb{P}(S_{100} = 0) = (1 - 0.03)^{100} \approx 0.0476.$$

The common- p assumption treats all impressions alike. We will revisit it when different users have different propensities to click.

2.6.3 Why rare Binomial successes look Poisson

The short-interval model contained many independent trials, each with a tiny success probability. A Binomial count behaves similarly whenever n is large and p is small, with $np = m$ held fixed. For any fixed k ,

$$\binom{n}{k} \left(\frac{m}{n}\right)^k \left(1 - \frac{m}{n}\right)^{n-k} \longrightarrow e^{-m} \frac{m^k}{k!}.$$

This is the Poisson limit derived earlier. Its moments agree with the Binomial moments in the limit:

$$np = m, \quad np(1-p) \longrightarrow m.$$

For the click example, the approximation is $\text{Poisson}(3)$, which gives no-click probability $e^{-3} \approx 0.0498$, close to the exact 0.0476.

2.7 Geometric waiting times

Now show the link repeatedly until the first click, instead of fixing the number of impressions in advance. Let G be the number of impressions needed, including the successful one. For example, the sequence

no click, no click, click

gives $G = 3$.

Assume independent trials with the same success probability $0 < p \leq 1$. For $0 < p < 1$, the event $G = k$ requires $k - 1$ failures followed by one success:

$$\mathbb{P}(G = k) = (1 - p)^{k-1}p, \quad k = 1, 2, \dots$$

This is the *Geometric distribution*, $G \sim \text{Geometric}(p)$. If $p = 1$, the first trial always succeeds and $G = 1$.

2.7.1 Waiting longer means accumulating failures

The event $G > k$ means that all of the first k trials failed. Hence

$$\mathbb{P}(G > k) = (1 - p)^k, \quad k = 0, 1, \dots$$

With $p = 0.1$, the chance of still waiting for a success after twenty trials is $0.9^{20} \approx 0.1216$.

This model has the same memoryless logic as exponential waiting. Once the first m trials have failed, the next trials still have the original success probability. Formally, for $0 < p < 1$,

$$\mathbb{P}(G > m + k \mid G > m) = \frac{(1 - p)^{m+k}}{(1 - p)^m} = (1 - p)^k.$$

An example is a scrolling session in which a user independently stops after each video with probability p . The number of videos watched is Geometric. If fatigue makes stopping more likely after each video, the constant- p assumption no longer holds.

2.7.2 Mean and variance

To compute the mean, weight each possible stopping trial by its probability. Write $q = 1 - p$. The geometric series

$$\sum_{k=0}^{\infty} q^k = \frac{1}{1 - q}, \quad |q| < 1,$$

can be differentiated to obtain the factor k needed in the mean:

$$\sum_{k=1}^{\infty} kq^{k-1} = \frac{1}{(1 - q)^2}.$$

Multiplying by p gives $\mathbb{E}[G] = 1/p$. A second differentiation gives

$$\mathbb{E}[G(G - 1)] = p \sum_{k=2}^{\infty} k(k - 1)q^{k-1} = \frac{2(1 - p)}{p^2}.$$

As in the Poisson calculation, use $G^2 = G(G - 1) + G$ to get

$$\text{Var}(G) = \frac{1 - p}{p^2}.$$

At $p = 0.1$, the mean is ten trials and the variance is ninety.

Some sources count failures before success instead. That variable is $H = G - 1$, so its mean is $(1 - p)/p$, its variance is unchanged, and it can equal zero. In these notes, Geometric always counts the successful trial too.

2.7.3 Turning a trial count into a time

Return to the inbox approximation, where one trial corresponds to a short interval of duration Δ minutes. A count of G intervals represents a waiting time of ΔG minutes. Set $p = \lambda\Delta$, just as in the first section.

By time t , we have completed $\lfloor t/\Delta \rfloor$ intervals; the floor symbol means the integer number of whole intervals. No arrival by that time has probability

$$\mathbb{P}(\Delta G > t) = (1 - \lambda\Delta)^{\lfloor t/\Delta \rfloor} \longrightarrow e^{-\lambda t} \quad (\Delta \downarrow 0).$$

The Geometric waiting count therefore becomes an Exponential waiting time as the time intervals shrink. The moments show the same connection:

$$\mathbb{E}[\Delta G] = \frac{1}{\lambda}, \quad \text{Var}(\Delta G) = \frac{1 - \lambda\Delta}{\lambda^2} \longrightarrow \frac{1}{\lambda^2}.$$

The time-unit conversion is what connects the two distributions: Geometric counts trials; Exponential measures elapsed time.

2.8 The Gaussian distribution

The waiting times studied so far are nonnegative and usually skewed: a wait can be much longer than average, but cannot extend below zero. Measurement errors pose a different problem. A sensor measuring a fixed quantity may sometimes read too high and sometimes too low. Several small disturbances can add together, with positive and negative contributions partly cancelling.

A useful model for such errors has a central peak and decreases symmetrically on either side. This is the *Gaussian*, also called the *normal*, distribution. It is often a useful approximation when many small independent disturbances contribute and no one disturbance dominates. The connection between sums and Gaussian approximations will be developed further in the chapter on limit theorems.

2.8.1 Center and width

Let μ describe the center of the errors and $\sigma > 0$ their scale. The center can be nonzero if the sensor is systematically biased. The density is

$$f_X(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left[-\frac{1}{2}\left(\frac{x - \mu}{\sigma}\right)^2\right], \quad x \in \mathbb{R}.$$

We write $X \sim \mathcal{N}(\mu, \sigma^2)$.

The expression $(x - \mu)/\sigma$ measures how far x lies from the center relative to the width of the curve. Squaring it treats positive and negative deviations equally. The negative exponential

makes values far from the center less likely. Increasing σ widens the curve; the factor $1/\sigma$ reduces its height to preserve total area one.

The simplest member is $Z \sim \mathcal{N}(0, 1)$, with density

$$\phi(z) = \frac{1}{\sqrt{2\pi}} e^{-z^2/2}.$$

The constant $\sqrt{2\pi}$ is the value of $\int_{-\infty}^{\infty} e^{-z^2/2} dz$, so this density integrates to one. Every other Gaussian can be obtained by shifting and scaling:

$$X = \mu + \sigma Z.$$

2.8.2 Why the parameters are the mean and standard deviation

Symmetry of ϕ means that the contributions of z and $-z$ to the mean cancel, giving $\mathbb{E}[Z] = 0$. For the second moment, differentiating the density gives $\phi'(z) = -z\phi(z)$. This turns the moment integral into an integration-by-parts calculation:

$$\begin{aligned} \mathbb{E}[Z^2] &= \int_{-\infty}^{\infty} z^2 \phi(z) dz = - \int_{-\infty}^{\infty} z \phi'(z) dz \\ &= [-z\phi(z)]_{-\infty}^{\infty} + \int_{-\infty}^{\infty} \phi(z) dz = 1. \end{aligned}$$

Thus $\text{Var}(Z) = 1$. Applying the shift and scale rules to $X = \mu + \sigma Z$ gives

$$\mathbb{E}[X] = \mu, \quad \text{Var}(X) = \sigma^2, \quad \text{sd}(X) = \sigma.$$

This justifies the parameters' interpretation. In $\mathcal{N}(\mu, \sigma^2)$, the second argument is the variance.

2.8.3 Expressing an error in standard-deviation units

For a sensor with mean error zero, an error of 0.4 units may be large or small relative to its usual variation. If the standard deviation is 0.2, that error is two standard deviations above the mean; if it is 2, it is only 0.2 standard deviations above the mean.

The transformation

$$Z = \frac{X - \mu}{\sigma}$$

subtracts the center and expresses the deviation in standard-deviation units. This is *standardization*. For a Gaussian X , the result always follows $\mathcal{N}(0, 1)$. We can therefore answer interval questions for any Gaussian using a single CDF:

$$\Phi(z) = \mathbb{P}(Z \leq z) = \int_{-\infty}^z \phi(u) du.$$

The endpoints of the original interval are transformed in the same way:

$$\mathbb{P}(a < X \leq b) = \Phi\left(\frac{b - \mu}{\sigma}\right) - \Phi\left(\frac{a - \mu}{\sigma}\right).$$

The integral defining Φ is evaluated numerically. Symmetry gives $\Phi(-z) = 1 - \Phi(z)$.

Example 2.10 (How often does a sensor make a large error?). Suppose the error is $E \sim \mathcal{N}(0, 0.2^2)$. An absolute error larger than 0.4 is an error more than two standard deviations from zero. The two symmetric tails give

$$\mathbb{P}(|E| > 0.4) = \mathbb{P}(|Z| > 2) = 2[1 - \Phi(2)] \approx 0.0455.$$

Under this model, about 4.55% of repeated readings have such an error. Likewise,

$$\mathbb{P}(|Z| \leq 1) \approx 0.6827, \quad \mathbb{P}(|Z| \leq 1.96) \approx 0.95.$$

The central 95% interval for the sensor error is therefore

$$[-1.96(0.2), 1.96(0.2)] = [-0.392, 0.392].$$

2.8.4 Adding errors and connecting back to earlier distributions

Changing units or adding a fixed bias preserves the Gaussian shape:

$$aX + b \sim \mathcal{N}(a\mu + b, a^2\sigma^2), \quad a \neq 0.$$

For $a = 0$, the result is simply the constant b .

Independent Gaussian errors also add to another Gaussian. If $X \sim \mathcal{N}(\mu_X, \sigma_X^2)$ and $Y \sim \mathcal{N}(\mu_Y, \sigma_Y^2)$ are independent, then

$$X + Y \sim \mathcal{N}(\mu_X + \mu_Y, \sigma_X^2 + \sigma_Y^2).$$

The convolution method used for Gamma waiting times gives this density. The means and variances agree with the addition rules we already know.

For example, two independent centered errors with standard deviations 0.2 and 0.3 produce a total error with variance $0.2^2 + 0.3^2 = 0.13$, and standard deviation $\sqrt{0.13} \approx 0.361$. Standard deviations do not simply add.

A sum can also be approximately Gaussian even when its individual terms are not Gaussian. We have already met two such sums: Binomial counts add Bernoulli outcomes, and integer-shape Gamma waits add exponential gaps. For sufficiently large parameters, useful approximations are

$$\begin{aligned} \text{Binomial}(n, p) &\approx \mathcal{N}(np, np(1-p)), \\ \text{Gamma}(k, \lambda) &\approx \mathcal{N}(k/\lambda, k/\lambda^2), \\ \text{Poisson}(m) &\approx \mathcal{N}(m, m). \end{aligned}$$

Each Gaussian uses the mean and variance of the distribution being approximated. For the Binomial, both the expected successes np and failures $n(1-p)$ should be large; for Gamma and Poisson, the approximations improve as k and m grow.

The Gaussian allows negative values, while counts and waiting times do not. Thus approximating a very small count or a strongly skewed waiting time by a Gaussian can assign substantial probability to impossible outcomes. The original distributions remain the appropriate starting point in those cases.

2.9 Mixtures and heterogeneity

Suppose two sensors measure the same reference quantity. One tends to read two units too low; the other tends to read two units too high. Within either sensor, the random error has standard deviation one.

Choose a sensor with equal probability for each reading. If we combine all readings without keeping track of which sensor produced them, what distribution do the errors follow? Averaging the two sensor biases gives zero, but that does not mean either sensor usually makes an error near zero. Most errors are still near -2 or 2 .

2.9.1 Choose a source, then draw an observation

Let Z record which sensor is selected:

$$\mathbb{P}(Z = 1) = \mathbb{P}(Z = 2) = \frac{1}{2}.$$

Given the chosen sensor, the error distribution is

$$X \mid Z = 1 \sim \mathcal{N}(-2, 1), \quad X \mid Z = 2 \sim \mathcal{N}(2, 1).$$

The notation $X \mid Z = k$ means that we are considering only readings from sensor k .

Write the two Gaussian densities as f_1 and f_2 . For any interval, half the readings come from each sensor, so the law of total probability gives half of each sensor's interval probability. The density is therefore

$$f_X(x) = \frac{1}{2}f_1(x) + \frac{1}{2}f_2(x).$$

This is a *mixture distribution*. Its two peaks reflect the two sources of readings.

The same construction works with K sources chosen with probabilities $\pi_1, \dots, \pi_K$:

$$f_X(x) = \sum_{k=1}^K \pi_k f_k(x), \quad \pi_k \geq 0, \quad \sum_{k=1}^K \pi_k = 1. \quad (2.6)$$

The f_k are called the mixture's *components*. For discrete observations, use component PMFs in the same formula. If the source label is not recorded, it is an unobserved, or *latent*, variable.

2.9.2 Where the pooled variance comes from

In the two-sensor example the overall mean is

$$\mathbb{E}[X] = \frac{1}{2}(-2) + \frac{1}{2}(2) = 0.$$

Within either sensor, the second moment equals its variance plus the square of its mean:

$$\mathbb{E}[X^2 \mid Z = 1] = 1 + (-2)^2 = 5, \quad \mathbb{E}[X^2 \mid Z = 2] = 1 + 2^2 = 5.$$

Averaging these gives $\mathbb{E}[X^2] = 5$. Hence the pooled variance is $5 - 0^2 = 5$, even though each sensor separately has variance one.

There are two sources of spread: random fluctuation within a sensor, and the difference between the two sensors' centers. To separate them in general, let component k have mean μ_k and variance σ_k^2 . Then

$$\mu = \mathbb{E}[X] = \sum_k \pi_k \mu_k, \quad \mathbb{E}[X^2] = \sum_k \pi_k (\sigma_k^2 + \mu_k^2).$$

Subtracting μ^2 and regrouping the terms gives

$$\boxed{\text{Var}(X) = \underbrace{\sum_k \pi_k \sigma_k^2}_{\text{variation within sources}} + \underbrace{\sum_k \pi_k (\mu_k - \mu)^2}_{\text{variation between sources}}.} \quad (2.7)$$

For the sensors, these contributions are 1 and 4, respectively.

The inner averages can be written as conditional expectations. The quantity $\mathbb{E}[X | Z]$ equals μ_k when source k is selected; it varies because the selected source varies. Similarly, $\text{Var}(X | Z)$ is the variance inside the selected source. The outer expectation averages over the choice of source. With this notation, the two identities become

$$\mathbb{E}[X] = \mathbb{E}[\mathbb{E}[X | Z]], \quad (2.8)$$

$$\text{Var}(X) = \mathbb{E}[\text{Var}(X | Z)] + \text{Var}(\mathbb{E}[X | Z]). \quad (2.9)$$

These are the *laws of total expectation and total variance*. We have derived them for a finite mixture with finite second moments.

2.9.3 The same mean and variance can hide different shapes

A single Gaussian $\mathcal{N}(0, 5)$ has the same mean and variance as the two-sensor mixture. Its most likely region is near zero, whereas the mixture has two peaks near the sensor biases (Figure 2.4).




Figure 2.4: Pooling two biased sensors produces the blue mixture density. The dotted curves are the weighted sensor densities, which add to the blue curve. A single Gaussian with the same mean and variance puts substantially more probability near zero.

For example, in sensor 1 the event $-1 \leq X \leq 1$ corresponds to a standardized error between one and three. Sensor 2 gives the same probability by symmetry. Therefore the pooled probability is

$$\mathbb{P}(|X| \leq 1) = \Phi(3) - \Phi(1) \approx 0.1573.$$

The Gaussian $\mathcal{N}(0, 5)$ instead gives $2\Phi(1/\sqrt{5}) - 1 \approx 0.3453$. Knowing the mean and variance alone does not determine interval probabilities.

Also distinguish selecting one sensor from averaging two readings. If we take one independent reading from *each* sensor and average them, their opposite biases cancel:

$$\frac{X_1 + X_2}{2} \sim \mathcal{N}(0, 1/2).$$

A mixture chooses which source supplies an observation. An average combines observations from both sources. These are different experiments.

2.9.4 Different users have different message rates

Two sensors give a small number of groups. In a user population, the differences can vary continuously. Some users receive very few messages; others receive many. This variation between members of a population is called *heterogeneity*.

Observe each user for one day. Let Λ be the user's expected number of messages in that day. A particular user's value, denoted λ , is fixed during the observation; the uppercase variable Λ varies when we choose a different user. Assume that, for a user with value λ ,

$$N \mid \Lambda = \lambda \sim \text{Poisson}(\lambda).$$

Thus even for a fixed user the realized count N fluctuates from day to day. Across users there is additional variation in the mean count.

The within-user mean and variance both equal Λ . Averaging over users, the laws of total expectation and variance give

$$\mathbb{E}[N] = \mathbb{E}[\Lambda], \quad \text{Var}(N) = \mathbb{E}[\Lambda] + \text{Var}(\Lambda).$$

The first contribution to the variance is daily Poisson variation. The second is variation in the users' average activity. If the latter is positive, the pooled variance exceeds the pooled mean. This is called *overdispersion* relative to a Poisson distribution.

Replacing every user's rate by the population average removes this second contribution. It also misses the extra zeros contributed by quiet users and the large counts contributed by busy users.

2.9.5 A Gamma distribution of Poisson means

We need a distribution for the positive quantity Λ . One possible choice is the Gamma family, whose shape and rate let us describe different amounts of variation among users:

$$\Lambda \sim \text{Gamma}(\alpha, \beta).$$

The Gamma moments immediately give

$$\mathbb{E}[N] = \frac{\alpha}{\beta}, \quad \text{Var}(N) = \frac{\alpha}{\beta} + \frac{\alpha}{\beta^2}.$$

For $\alpha = 2, \beta = 1/2$, the mean user count is four and the variance between user means is eight. The pooled count therefore has mean four and variance $4 + 8 = 12$.

We can also find the full count distribution. Let $g(\lambda)$ be the Gamma density of user means. To find the probability of k messages, average the Poisson probability over all possible user means:

$$\mathbb{P}(N = k) = \int_0^\infty \mathbb{P}(N = k \mid \Lambda = \lambda) g(\lambda) d\lambda.$$

This is the continuous version of the weighted sum for the sensors: possible user means replace the discrete source labels, and the integral replaces the sum. The count remains discrete; it is the mixing parameter that varies continuously.

Substituting the Poisson PMF and Gamma density gives

$$\begin{aligned} \mathbb{P}(N = k) &= \int_0^\infty \frac{e^{-\lambda} \lambda^k}{k!} \frac{\beta^\alpha}{\Gamma(\alpha)} \lambda^{\alpha-1} e^{-\beta\lambda} d\lambda \\ &= \frac{\Gamma(k + \alpha)}{\Gamma(\alpha) k!} \left(\frac{\beta}{\beta + 1} \right)^\alpha \left(\frac{1}{\beta + 1} \right)^k, \quad k = 0, 1, \dots \end{aligned}$$

The second line follows by combining the powers and exponentials, then using the gamma integral with shape $k + \alpha$ and rate $\beta + 1$. The resulting distribution is called *negative binomial*. Here its meaning is supplied by the construction: a Poisson count whose mean differs from user to user according to a Gamma distribution.

For our numerical example,

$$\mathbb{P}(N = 0) = \left(\frac{1/2}{1 + 1/2} \right)^2 = \frac{1}{9} \approx 0.1111.$$

A single $\text{Poisson}(4)$ distribution instead gives $e^{-4} \approx 0.0183$. The mean counts agree, but the predicted fraction of users receiving no messages is very different.




Figure 2.5: Both populations average four messages per user per day. With a common Poisson mean of four, the variance is four. With Gamma-distributed user means, the variance is twelve, with more users receiving zero messages or unusually many messages.

The conditional-moment rules used here extend the finite-mixture identities by replacing weighted sums with integrals. They require finite moments, but they do not require us to evaluate the full mixture PMF before calculating its mean and variance.

2.9.6 Ten responses from one user or one response from ten users?

Heterogeneity also matters for binary observations, but the way we group observations changes the result. Suppose half the users click with probability 0.1, and half with probability 0.9. A randomly chosen user's response is a click with probability

$$\frac{1}{2}(0.1) + \frac{1}{2}(0.9) = 0.5.$$

That overall probability does not tell us how much the individual user probabilities differ.

Now compare two experiments. In the first, choose one user and show them ten posts, with independent responses once their click probability is fixed. The chosen user is likely to click about once or about nine times, depending on which group they belong to. In the second, choose ten users independently and record one response from each. Each response is independently a click with probability 0.5, so the count is $\text{Binomial}(10, 0.5)$.

The means are both five. The variances are different. Write P for the first experiment's shared user probability and S for the ten-response count. Within either user group, the variance is

$10(0.1)(0.9) = 0.9$. The group means are one and nine, whose variance around five is sixteen. Therefore

$$\text{Var}(S) = 0.9 + 16 = 16.9,$$

whereas the second experiment has variance $10(0.5)(0.5) = 2.5$.

More generally, let $q = \mathbb{E}[P]$ and suppose $S \mid P \sim \text{Binomial}(n, P)$. Then

$$\begin{aligned}\mathbb{E}[S] &= nq, \\ \text{Var}(S) &= n\mathbb{E}[P(1-P)] + n^2 \text{Var}(P) \\ &= nq(1-q) + n(n-1) \text{Var}(P).\end{aligned}$$

The final step uses $\mathbb{E}[P(1-P)] = q(1-q) - \text{Var}(P)$, obtained by expanding $\mathbb{E}[P^2] = \text{Var}(P) + q^2$. The extra variance comes from sharing the same random user probability across trials.

If a fresh P_i is independently drawn for each trial instead, and the responses are independently generated given those probabilities, the responses are independent Bernoulli variables with parameter q . Their sum is $\text{Binomial}(n, q)$. Which quantities are shared across observations is part of specifying the experiment.

2.10 Extension: Gaussian scale mixtures and Student's t

Return to measurement errors. So far the two sensors differed in their average errors. Another possibility is that a sensor is usually precise but occasionally operates under much noisier conditions, without changing its average error.

For example, suppose 95% of readings have error $\mathcal{N}(\mu, \sigma^2)$, and the remaining 5% have standard deviation 5σ :

$$X \sim 0.95 \mathcal{N}(\mu, \sigma^2) + 0.05 \mathcal{N}(\mu, 25\sigma^2).$$

This is called a *contaminated Gaussian* model. The noisier readings make large errors more common. The mean stays μ , while the mixture variance is

$$\text{Var}(X) = 0.95\sigma^2 + 0.05(25\sigma^2) = 2.2\sigma^2.$$

Instead of two noise levels, we can allow a continuous range. Describe the noise using its inverse variance τ , often called *precision*. A small τ means a large variance and a noisy reading; a large τ means a narrow error distribution. For a standardized error, use

$$Z \mid \tau \sim \mathcal{N}(0, 1/\tau).$$

The Gamma family is a natural candidate for a positive random precision. Choose

$$\tau \sim \text{Gamma}(\nu/2, \nu/2), \quad \nu > 0.$$

Using equal shape and rate fixes $\mathbb{E}[\tau] = 1$; the Gamma variance formula gives $\text{Var}(\tau) = 2/\nu$. Thus ν controls how much the noise level varies: large ν keeps precision close to one, while small ν permits more variation.

Average the conditional Gaussian density over these precision values, just as we averaged Poisson probabilities over user means. Let g_τ denote the Gamma density and $f_{Z|\tau}(z \mid u)$ the

Gaussian density with variance $1/u$. Then

$$\begin{aligned} f_Z(z) &= \int_0^\infty f_{Z|\tau}(z | u) g_\tau(u) du \\ &= \frac{\Gamma((\nu+1)/2)}{\sqrt{\nu\pi} \Gamma(\nu/2)} \left(1 + \frac{z^2}{\nu}\right)^{-(\nu+1)/2}. \end{aligned}$$

The product inside the integral is proportional, as a function of u , to $u^{(\nu-1)/2} e^{-(\nu+z^2)u/2}$. This is another gamma integral, now with shape $(\nu+1)/2$ and rate $(\nu+z^2)/2$. The resulting density defines *Student's t -distribution*, written $Z \sim t_\nu$. In this construction, ν controls variation in precision and therefore how often extreme errors occur.

Occasionally drawing a small precision produces an unusually large Gaussian error. After averaging over precisions, extreme errors remain much more likely than under a standard Gaussian. This is the meaning of *heavy tails* here: for large $|z|$, the density decreases like $|z|^{-(\nu+1)}$, rather than like $e^{-z^2/2}$.

For $\nu > 1$, the mean exists and is zero. For $\nu > 2$, total variance and the Gamma integral give

$$\text{Var}(Z) = \mathbb{E}[1/\tau] = \frac{\nu}{\nu-2}.$$

When $1 < \nu \leq 2$, rare large observations make the second-moment integral diverge, so the variance is infinite. When $0 < \nu \leq 1$, even the absolute first moment diverges and the mean is undefined. This illustrates why the convergence conditions in the definitions of expectation and variance matter.

For an error with center μ and scale $s > 0$, write $X = \mu + sZ$. For $\nu > 2$,

$$\mathbb{E}[X] = \mu, \quad \text{Var}(X) = s^2 \frac{\nu}{\nu-2}.$$

The scale s is therefore not its standard deviation. As ν increases, precision concentrates near one and the standard Student distribution approaches $\mathcal{N}(0, 1)$.

2.11 Distribution reference

The table collects the means and variances derived above. The support lists the values the variable can take, distinguishing counts from waiting times and signed errors.

Distribution	Support	$\mathbb{E}[X]$	$\text{Var}(X)$
$\text{Exp}(\lambda)$	$[0, \infty)$	$1/\lambda$	$1/\lambda^2$
$\text{Poisson}(m)$	$0, 1, \dots$	m	m
$\text{Gamma}(\alpha, \beta)$	$[0, \infty)$	α/β	α/β^2
$\text{Bernoulli}(p)$	$0, 1$	p	$p(1-p)$
$\text{Binomial}(n, p)$	$0, \dots, n$	np	$np(1-p)$
$\text{Geometric}(p)$	$1, 2, \dots$	$1/p$	$(1-p)/p^2$
$\mathcal{N}(\mu, \sigma^2)$	$\mathbb{R}$	μ	σ^2
t_ν	$\mathbb{R}$	$0 \ (\nu > 1)$	$\nu/(\nu-2) \ (\nu > 2)$

Here $\lambda, \alpha, \beta, m, \nu > 0$, $\sigma > 0$, and n is a nonnegative integer. The Bernoulli and Binomial parameters allow $0 \leq p \leq 1$; the Geometric requires $0 < p \leq 1$. Gamma uses shape and rate. Geometric counts trials including the first success. Student moments exist only under the conditions shown.

Connection	Required mechanism or condition
Exponential to Gamma	Sum independent waiting times with a common rate.
Exponential to Poisson	Count arrivals generated by independent exponential interarrival times.
Bernoulli to Binomial	Sum a fixed number of independent trials with the same success probability.
Binomial to Poisson	Take $n \rightarrow \infty$, $p \rightarrow 0$, with $np = m$.
Geometric to Exponential	Rescale trial duration to Δ , set $p = \lambda\Delta$, and let $\Delta \downarrow 0$.
Poisson to negative binomial	Mix the Poisson rate with a Gamma distribution.
Gaussian to Student	Mix the Gaussian precision with $\text{Gamma}(\nu/2, \nu/2)$.